

# EpiFusion Tutorial

2023-12-20

## Introduction

This R Markdown document will guide you through the process of parsing and interpreting the output of the EpiFusion tutorial. We will use a selection of functions that can also be found in the **EpiFusion\_utilities.R** script on the GitHub repository.

The first step of the tutorial is to point the document to the filepath of the output folder made by EpiFusion. We will also source the **EpiFusion\_utilities.R** script to load the necessary functions to look through this output. The current filepath assumes the folder is in your working directory; you may need to change this depending on where your output folder was generated.

```
outputfolder <- "intro-example/"
source("EpiFusion_utilities.R")
```

```
## Loading required package: pacman
```


```
## [1] "Loading EpiFusion functions."
```

NOTE: the first time you run this, the output might also show the installation of any required R packages you don't already have installed.

## Checking convergence

The first thing we'll do is take a look at the trace plots of the chain likelihood, to check that the model has run nicely. For this we can use the **plotlikelihoodtrace** function, which takes the output folder filepath as its input.

## Log Likelihood Trace (full chain)



Have a look at the likelihood trace for each chain run. They should start in different places but eventually converge to roughly the same value. Keep an eye out for chains getting stuck (staying at the same value for a long time). This does happen occasionally, and we are working on introducing adaptive Metropolis Hastings MCMC sampling to EpiFusion to fix this. Once you have had a look at the trace plot you can choose what proportion of each chain to discard as burn-in. EpiFusion models tend to converge relatively quickly, so the default in this document will be 10%. However, you can edit this as needed.

```
burn_in <- 0.1
```

## Examining infection trajectory posteriors

### Trajectories table

Next we'll examine the infection trajectories inferred by your EpiFusion model. First let's take a look at the raw values by loading the trajectories to our environment with `loadtrajectoriesminusburnin`, which takes the `outputfolder` and `burn_in` as arguments. This will return a table, with a column for every day of the analysis and a row for every sampled trajectory. Below we load the table and look at the first sample.

```
trajectories <- loadtrajectoriesminusburnin(outputfolder, burn_in)
unlist(trajectories[1,])
```

##	T_0	T_1	T_2	T_3	T_4	T_5	T_6	T_7	T_8	T_9	T_10	T_11	T_12
##	1	1	1	4	6	7	7	10	15	21	25	30	35
##	T_13	T_14	T_15	T_16	T_17	T_18	T_19	T_20	T_21	T_22	T_23	T_24	T_25
##	40	47	53	65	67	73	83	89	113	126	146	168	205
##	T_26	T_27	T_28	T_29	T_30	T_31	T_32	T_33	T_34	T_35	T_36	T_37	T_38
##	227	271	309	350	396	452	515	557	634	705	772	833	894

```
## T_39 T_40 T_41 T_42 T_43 T_44 T_45 T_46 T_47 T_48 T_49 T_50 T_51
## 987 1066 1195 1294 1437 1611 1796 1987 2261 2589 2900 3210 3509
## T_52 T_53 T_54 T_55 T_56 T_57 T_58 T_59 T_60 T_61 T_62 T_63 T_64
## 3765 3906 4063 4159 4154 4182 4167 4146 4060 3946 3781 3605 3470
## T_65 T_66 T_67 T_68 T_69 T_70 T_71 T_72 T_73 T_74 T_75 T_76 T_77
## 3305 3174 3000 2842 2726 2560 2453 2289 2092 1886 1685 1552 1440
## T_78 T_79 T_80 T_81 T_82 T_83 T_84 T_85 T_86 T_87 T_88 T_89 T_90
## 1343 1225 1109 1001 910 832 757 690 632 575 537 514 474
## T_91 T_92 T_93 T_94 T_95 T_96 T_97 T_98 T_99 T_100 T_101 T_102 T_103
## 429 395 372 333 312 293 266 241 224 209 182 161 152
## T_104 T_105 T_106 T_107 T_108 T_109 T_110 T_111 T_112 T_113 T_114 T_115 T_116
## 145 136 125 117 106 102 95 88 80 75 65 61 50
## T_117 T_118 T_119 T_120 T_121 T_122 T_123 T_124 T_125 T_126 T_127 T_128 T_129
## 44 41 39 35 32 35 35 39 36 33 28 28 25
## T_130 T_131 T_132 T_133 T_134 T_135 T_136 T_137 T_138 T_139 T_140
## 25 22 19 18 17 17 15 15 15 15 15
```

You can also load the trajectories inferred by each chain separately. This can be useful if you think you might have convergence issues.


```
trajectories_separate <- loadtrajectoriesminusburninseparate(outputfolder, burn_in)
```

This function returns a list of length (N) of trajectory tables, where N = the number of chains.

## Trajectories plot

It can be difficult to get a sense for what the trajectories are really saying in table form, so let's plot them using `plottrajectoriesfromtable`. We can supply our desired colour as the second argument to the function. There is also a function, `plottrajectoryposteriors`, which can plot the trajectories directly from the file, but we won't use that here. If you do want to use it, it takes `outputfolder`, `burn_in` and `colour` as arguments.

```
plottrajectoriesfromtable(trajectories, 'slateblue')
```



This function plots the mean sampled trajectory as a coloured line, with 95%, 80% and 66% HPDs as shaded regions of increasing darkness.

## Examining $R(t)$ trajectory posteriors

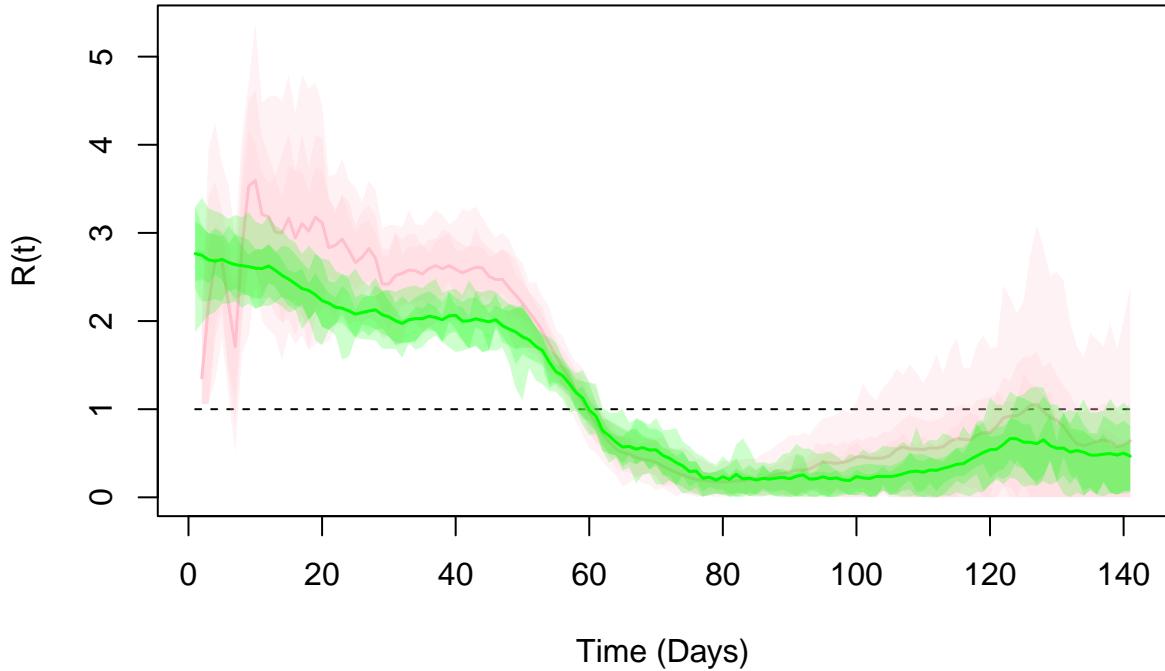
Now let's take a look at the  $R(t)$  estimates from the model. With EpiFusion it is actually possible to get  $R(t)$  through two methods. We'll examine both here.

### Renewal equation or Beta/Gamma?

The EpiFusion program stores daily infection events in the trajectories that it models, and uses this along with a user specified probability mass function of the generation time in a renewal equation to calculate an  $R(t)$  trajectory for each MCMC sample. These are stored in the `rt_chainN.csv` files of the EpiFusion output. However, EpiFusion also samples trajectories of the force of infection beta over time (`betas_chainN.csv`). Together with the posterior samples for the recovery/removal parameter gamma, you can also get  $R(t)$  estimates from dividing  $\text{beta}(t)/\text{gamma}(t)$  for each MCMC sample. Let's plot both options together to compare:

```
plotrtposteriors_renewal(outputfolder, burn_in, 'pink')
addrtposteriors_betagamma(outputfolder, burn_in, 'green')
```

## R(t) Trajectories



Here we've plotted the renewal  $R(t)$  posteriors with `plotrtposteriors_renewal` and added a layer of the beta/gamma  $R(t)$  posteriors with `addrtposteriors_betagamma`. There is also the inverse of each of these functions if you wanted to reverse the order; i.e. `addrtposteriors_renewal` and `plotrtposteriors_betagamma`. These  $R(t)$  trajectories should be in agreement, however the renewal equation version can be a little unstable when prevalence is low (at the start and end of this example). You can choose whichever option meets your needs.

### Loading $R(t)$ tables

You can load the  $R(t)$  in table form also, as shown below. The first value of the renewal equation trajectories will always be NaN due to the edge limitations of this method.

```
#Load renewal equation trajectories
renewal_rts <- loadrtsminusburnin(outputfolder, burn_in)
#Load beta/gamma trajectories
betagamma_rts <- loadbetagammartminusburnin(outputfolder, burn_in)
```

Similarly to the infection trajectories, there are also functions for plotting  $R(t)$  posteriors (or adding layers) directly from the tables. A list of all available functions is on the GitHub wiki, we recommend going through it to see what's available.

### MCMC Parameter Posteriors

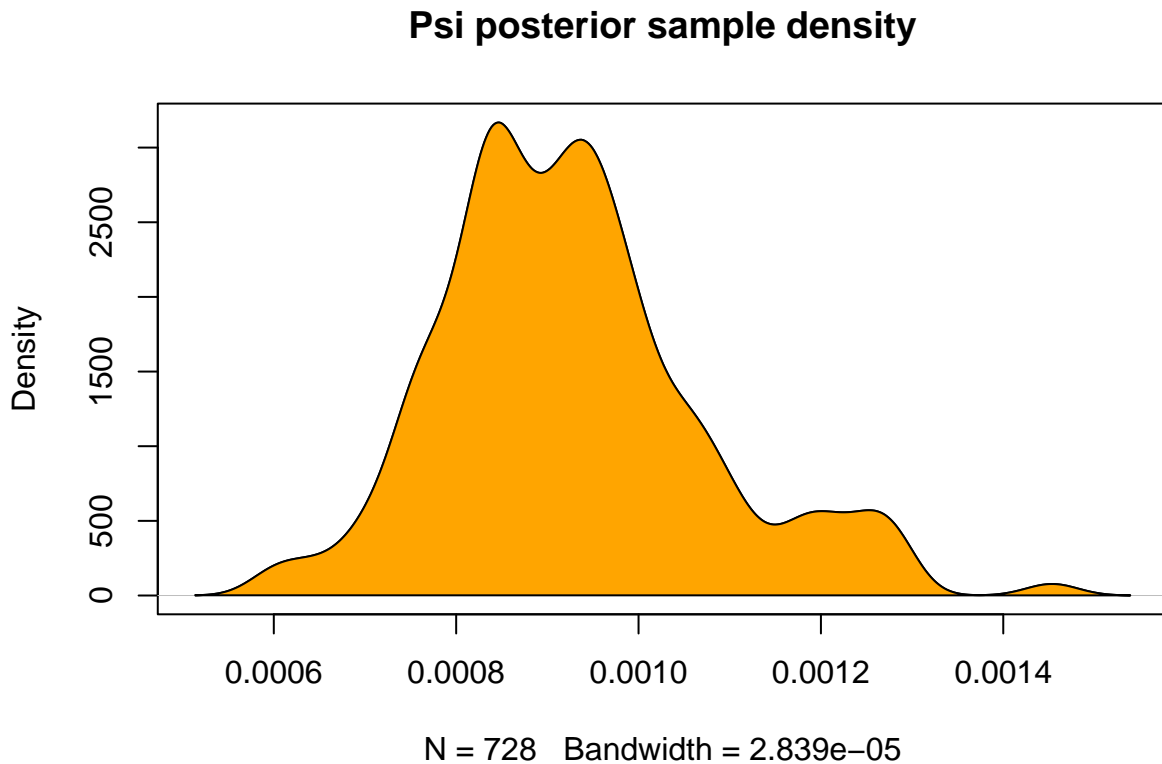
Another set of files you'll find in the EpiFusion are the `params_chainN.txt` files. These contain a row for each MCMC sample, and a column for each MCMC parameter. You can load them using the `loadparamsminusburnin` function. The last column of the table will always contain NAs, this is just an artefact of how EpiFusion saves the samples.

```
paramposteriors <- loadparamsminusburnin(outputfolder, burn_in)
head(paramposteriors)
```

```
##      gamma      psi      phi initialBeta betaJitter  X
## 20 0.1015715 0.0008628258 0.01966505  0.2898188 0.01303787 NA
## 21 0.1043629 0.0007787591 0.01939411  0.2795766 0.01628408 NA
## 22 0.1144166 0.0008406108 0.01896374  0.2722224 0.01703538 NA
## 23 0.1144166 0.0008406108 0.01896374  0.2722224 0.01703538 NA
## 24 0.1204064 0.0008672240 0.01728061  0.3067304 0.01781600 NA
## 25 0.1204064 0.0008672240 0.01728061  0.3067304 0.01781600 NA
```

If you are interested in a specific parameter, for example 'psi', you can load the posterior samples for that specific parameter as a vector with `loadparambyname`.

```
psiposterior <- loadparambyname(outputfolder, burn_in, 'psi')
plot(density(psiposterior), main = 'Psi posterior sample density')
polygon(density(psiposterior), col = 'orange')
```



## Conclusion

The above covers the basics of parsing EpiFusion output, and introduces some of the functions of `EpiFusion_utilities.R`. We stress that EpiFusion is still in its infancy, with more improvements happening every day, but we hope this tutorial will have peaked your interest! For more information we recommend checking out the EpiFusion wiki.